

Learning Continuous Control Policies by Stochastic Value Gradients (SVG) (1510.)

- Background
 - discrete-time MDP with continuous states and actions
 - time-varying reward function
 - either finite-horizon (time-dependent value functions) or infinite-horizon (stationary value functions) sum of rewards
- Stochastic value gradients
 - Backpropagation in stochastic Bellman equation.
 - Compute gradients along real trajectories \rightarrow SVG(∞).
 - Integrate value function critics \rightarrow SVG(1), extending to infinite-horizon control.
 - In the special case of additive noise, derivatives of model are noise independent. Otherwise, use re-parameterized generative model to infer the missing noise variables.
- SVG(∞)
 - Algorithm:
 - repeat
 - sample one length- T trajectory under π , append to \mathcal{D}
 - train \hat{f} with \mathcal{D}
 - on the latest trajectory:
 - infer $\xi|(s, a, s')$ and $\eta|(s, a)$
 - do backward recursions from T down to 0:

$$v_\theta = \left[r_a \pi_\theta + \gamma \left(v'_{s'} \hat{f}_a \pi_\theta + v'_\theta \right) \right] \Big|_{\eta, \xi}$$

$$v_s = \left[r_s + r_a \pi_s + \gamma v'_{s'} \left(\hat{f}_s + \hat{f}_a \pi_s \right) \right] \Big|_{\eta, \xi}$$
 - update policy with gradient v_θ^0
 - On-policy.
- SVG(1)-ER
 - Algorithm:
 - repeat
 - generate transitions under π_θ , append to \mathcal{D}
 - train \hat{f} on \mathcal{D}
 - train \hat{V} on \mathcal{D}
 - sample transitions from \mathcal{D} generated in previous loops
 - compute empirically estimated importance-weighting gradient (which involves value function)

- update policy with the gradient
- Critic reduces the variance of the gradient estimates.
- Extends to infinite-horizon.
- Off-policy.

Neural Network Dynamics for Model-Based Deep Reinforcement Learning with Model-Free Fine-Tuning (1708.)

- Model-based RL
 - Dynamics function predicts state change over Δt .
 - Training the learned dynamics function:
 - collecting training data with random exploration, off-policy
 - data preprocessing: normalization and Gaussian corruption
 - train the model with H -step validation errors (multi-step open-loop predictions)
 - random-sampling shooting with MPC, use model to estimate optimal action sequence and execute the first action
 - Algorithm:
 - Gather random trajectories ($\mathcal{D}_{\text{RAND}}$)
 - repeat
 - update \hat{f}_θ with gradient descent on MSE ($\mathcal{D}_{\text{RAND}}$ and \mathcal{D}_{RL})
 - model-based MPC controller gathers T new on-policy (real) transitions (\mathcal{D}_{RL})
- MB-MF:
 - MF learner initialization
 - gather example trajectories with MPC controller with pre-learned \hat{f}
 - train a neural network Gaussian policy with learned mean and fixed covariance to match expert trajectories via SGD on MSE
 - apply DAGGER
 - MF RL
 - TRPO

Model-Based Reinforcement Learning via Meta-Policy Optimization (MB-MPO) (1809.)

- Meta-RL learns an initialization θ^* such that for any task $\mathcal{M}_k \sim \rho(\mathcal{M})$ the policy attains maximum performance in the respective task after one policy gradient step.

$$\max_{\theta} \mathbb{E}_{\mathcal{M}_k \sim \rho(\mathcal{M}), f_k, \pi_{\theta'}} \left[\sum_{t=0}^{H-1} r_k(s_t, a_t) \right] \quad \text{s.t.: } \theta' = \theta + \alpha \nabla_{\theta} \mathbb{E}_{f_k, \pi_{\theta}} \left[\sum_{t=0}^{H-1} r_k(s_t, a_t) \right]$$

- MB-MPO
 - Model learning: models are decorrelated with random initialization and different subsets of random samples; dynamic models retrained via MPC controller with warm starts

- Meta-RL on learned models
 - meta-objective:

$$\max_{\theta} \frac{1}{K} \sum_{k=0}^K J_k(\theta'_k) \quad \text{s.t.:} \quad \theta'_k = \theta + \alpha \nabla_{\theta} J_k(\theta)$$
 where $J_k(\theta) := \mathbb{E}_{\mathbf{a}_t \sim \pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)} \left[\sum_{t=0}^{H-1} r(\mathbf{s}_t, \mathbf{a}_t) \mid \mathbf{s}_{t+1} = \hat{f}_{\phi_k}(\mathbf{s}_t, \mathbf{a}_t) \right]$
 - use imaginary trajectories, off-policy
 - TPRO for maximizing meta-objective, VPG for adaptation objectives $J_k(\theta)$
- Algorithm:
 - initialize policy π_{θ} , models $\hat{f}_{\phi_1}, \hat{f}_{\phi_2}, \dots, \hat{f}_{\phi_K}$
 - repeat
 - sample real trajectories with adapted policies $\pi_{\theta'_1}, \dots, \pi_{\theta'_K}$, append to \mathcal{D}
 - train all models with \mathcal{D}
 - for all models do
 - sample imaginary trajectories \mathcal{T}_k from \hat{f}_{ϕ_k} using π_{θ}
 - update adapted policy with \mathcal{T}_k
 - sample imaginary trajectories \mathcal{T}'_k from \hat{f}_{ϕ_k} using $\pi'_{\theta'_k}$
 - update θ with $\theta \rightarrow \theta - \beta \frac{1}{K} \sum_k \nabla_{\theta} J_k(\theta'_k)$ using \mathcal{T}'_k

Proximal Policy Optimization Algorithms (PPO) (1707.)

- Proposed objective function enables multiple epochs of minibatch updates.
- Background
 - Policy Gradient
 - estimator $\hat{g} = \hat{\mathbb{E}}_t \left[\nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t) \hat{A}_t \right]$
 - Trust Region
 - maximize $\hat{\mathbb{E}}_t \left[\frac{\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)}{\pi_{\theta_{\text{old}}}(\mathbf{a}_t | \mathbf{s}_t)} \hat{A}_t \right]$
 - subject to $\hat{\mathbb{E}}_t [\text{KL}[\pi_{\theta_{\text{old}}}(\cdot | \mathbf{s}_t), \pi_{\theta}(\cdot | \mathbf{s}_t)]] \leq \delta$
 - approximately solved using conjugate gradient algorithm after linear approximation to objective and quadratic approximation to constraint
- Surrogate objective
 - clipped surrogate objective $L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min \left(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right]$
 - adaptive KL penalty coefficient
 - $L^{KL PEN}(\theta) = \hat{\mathbb{E}}_t \left[\frac{\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)}{\pi_{\theta_{\text{old}}}(\mathbf{a}_t | \mathbf{s}_t)} \hat{A}_t - \beta \text{KL}[\pi_{\theta_{\text{old}}}(\cdot | \mathbf{s}_t), \pi_{\theta}(\cdot | \mathbf{s}_t)] \right]$
 - optimized via several epochs of minibatch SGD
 - $d = \hat{\mathbb{E}}_t [\text{KL}[\pi_{\theta_{\text{old}}}(\cdot | \mathbf{s}_t), \pi_{\theta}(\cdot | \mathbf{s}_t)]]$

if $d < d_{\text{targ}} / 1.5$, $\beta \leftarrow \beta / 2$; if $d > d_{\text{targ}} / 1.5$, $\beta \leftarrow \beta \times 2$

- If policy and value function share parameters, maximize CLIP + VF + S where CLIP = clipped surrogate objective, VF = squared-error loss for value function, S = entropy term to encourage exploration.
- Proximal Policy Optimization (PPO), AC Style
 - Algorithm:
 - repeat
 - for all N actors do
 - collect on-policy real transitions for T timesteps
 - compute advantage estimate $\hat{A}_1, \dots, \hat{A}_T$ (to compute loss)
 - update policy via gradients from surrogate loss L , with K epochs and minibatch size $M \leq NT$

Model-Ensemble Trust-Region Policy Optimization (ME-TRPO) (1802.)

- Vanilla Model-Based Deep RL
 - Algorithm:
 - repeat
 - collect real samples, add to \mathcal{D}
 - train \hat{f} with \mathcal{D}
 - repeat
 - collect fictitious samples with \hat{f}
 - BPTT
 - estimate expected sum of rewards
- Model Ensemble Trust Region Policy Optimization (ME-TRPO)
 - Algorithm:
 - repeat
 - collect real samples, add to \mathcal{D}
 - train all K models with \mathcal{D}
 - for each model in the ensemble (in parallel) # optimize policy using all models
 - collect fictitious samples with model
 - update policy using TRPO
 - estimate expected sum of rewards

Model-Based Value Expansion for Efficient Model-Free Reinforcement Learning (MVE) (1803)

- MVE is a value estimate for a policy π by assuming the model is accurate in depth H

- One sufficient condition under which MVE MSE \leq underlying critic MSE: given an arbitrary off-policy distribution of state-action pairs β , set $\nu = \mathbb{E} \left[(f^\pi)^T \beta \right]$ where $T \sim \text{Uniform} \{0, \dots, H-1\}$.
- $\text{MSE}_\nu(\hat{V}) \rightarrow$ Bellman error w.r.t. $\nu \rightarrow$ using target $\hat{V}_k(\hat{s}_T)$ which is equivalent to train \hat{V} with imaginary TD- k error. The target is used to train \hat{V} on the entire support of ν instead of just β . Sampling transitions from ν is equivalent to sampling from any point up to H imagined steps into the future, when starting from a state sampled from β .
- Algorithm: MVE
 - initialize targets $\theta' = \theta$ for π , $\varphi' = \varphi$ for Q
 - initialize replay buffer $\beta \leftarrow \emptyset$
 - repeat
 - collect transitions, add to β (empirical distribution of transitions observed in real world)
 - fit dynamics \hat{f} with LSE
 - repeat
 - sample $\tau_0 \sim \beta$
 - update θ with $\nabla_{\theta} \ell_{\text{actor}}(\pi_{\theta}, Q_{\varphi}, \tau_0)$
 - imagine future transitions τ_t up to H
 - update φ with $\nabla_{\varphi} \sum_t \ell_{\text{critic}}^{\pi_{\theta'}, \hat{Q}_{H-t}}(\varphi, \tau_t) / H$ # ν -based Bellman error of Q_{φ} proxied by k -step MVE of $Q_{\varphi'}$
 - update targets θ', φ' with some decay

Algorithmic Framework For Model-based Deep RL with Theoretical Guarantees (1807.)

- Meta-Algorithm for MB RL with some designed discrepancy bound D and distance function d
 - For $k = 0 \dots T$

$$\pi_{k+1}, M_{k+1} = \underset{\pi \in \Pi, M \in \mathcal{M}}{\operatorname{argmax}} V^{\pi, M} - D_{\pi_k, \delta}(M, \pi)$$

$$\text{s.t. } d(\pi, \pi_k) \leq \delta$$
 - It iteratively optimizes the lower bound over the policy π_{k+1} and the model M_{k+1} , subject to the constraint that the policy is not very far from the reference policy π_k obtained in the previous iteration.
 - The policy performance in the real environment is non-decreasing under the assumption that the real dynamics belongs to the parameterized family \mathcal{M}
- The meta-algorithm is instantiated by Stochastic Lower Bound Optimization (SLBO)
 - Algorithm:
 - repeat

- repeat
 - collect real samples, add to \mathcal{D}
 - repeat
 - repeat
 - update \hat{f} with SGD on H -step loss
 - repeat
 - collect fictitious samples under \hat{f} as \mathcal{D}'
 - optimize π_θ on \mathcal{D}' by TRPO with entropy regularization

Model-Predictive Policy Learning with Uncertainty Regularization for Driving in Dense Traffic (1901.)

- Model-predictive policy learning with uncertainty regularization (MPUR)
 - action-conditional forward model
 - per-sample loss:

$$\mathcal{L}(\theta, \phi; s_{1:t}, s_{t+1}, a_t) = \|s_{t+1} - f_\theta(s_{1:t}, a_t, z_t)\|_2^2 + \beta D_{KL}(q_\phi(z|s_{1:t}, s_{t+1}) \| p(z))$$
 - apply z -dropout to decouple latent variables and variation in prediction model outputs due to actions
 - training policy network with uncertainty minimization (MPUR) / expert regularization (MPER)
 - use forward model to train policy network π_ψ
 - sample initial state sequence $s_{1:t}$ from the training set
 - unroll forward model over T time steps
 - backpropagation w.r.t. ψ
 - objective: minimize policy cost + uncertainty cost

$$\underset{\psi}{\operatorname{argmin}} \left[\sum_{i=1}^T C(\hat{s}_{t+i}) + \lambda U(\hat{s}_{t+i}) \right], \text{ such that: } \begin{cases} z_{t+i} \sim p(z) \\ \hat{a}_{t+i} \sim \pi_\psi(\hat{s}_{t+i-1}) \\ \hat{s}_{t+i} = f(\hat{s}_{t+i-1}, \hat{a}_{t+i}, z_{t+i}) \end{cases}$$